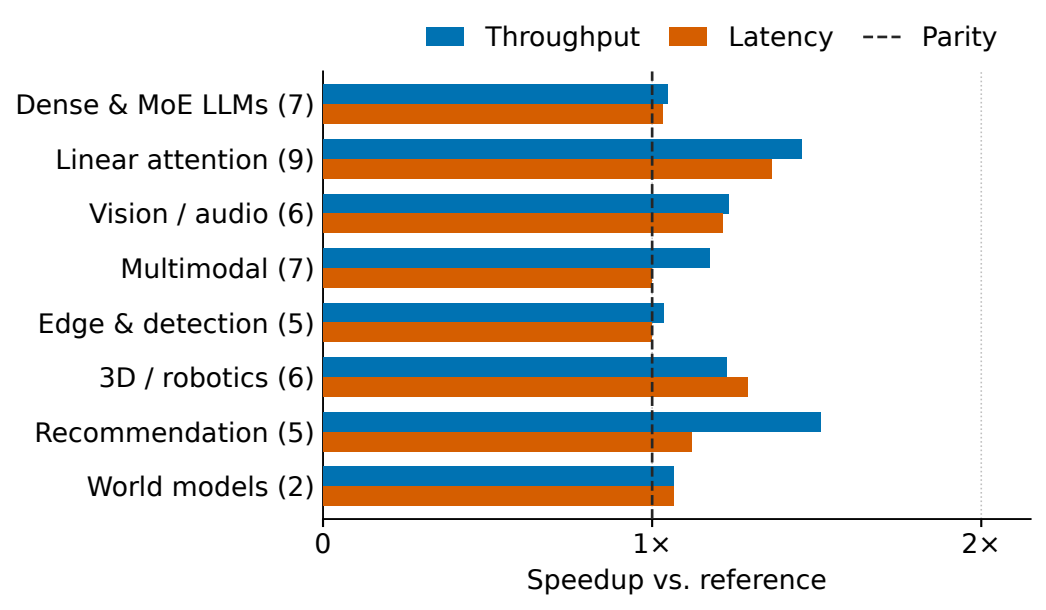
Throughput
Latency
Parity
Dense & MoE LLMs (7)
Linear attention (9)
Vision / audio (6)
Multimodal (7)
Edge & detection (5)
3D / robotics (6)
Recommendation (5)
World models (2)
0
1×
2×
Speedup vs. reference