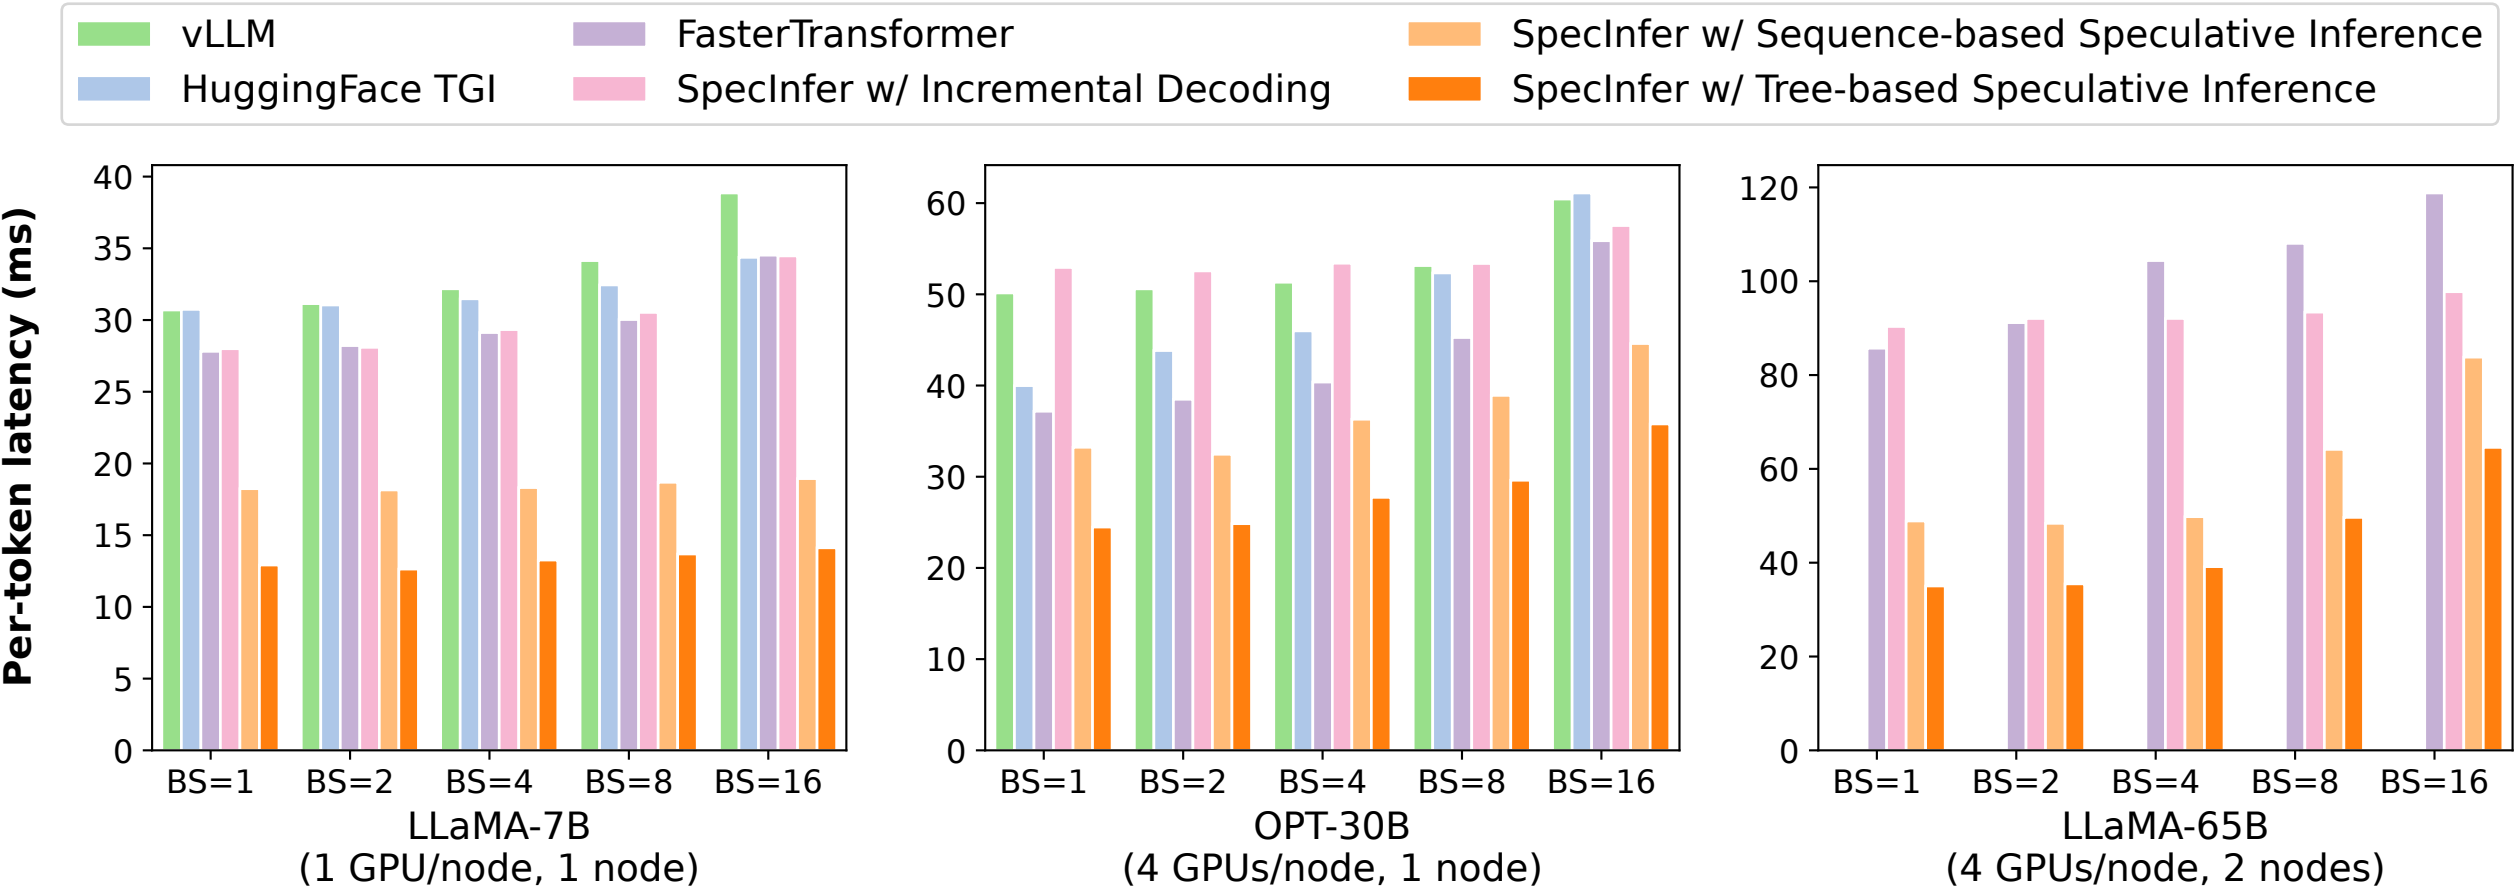
vLLM
HuggingFace TGI
FasterTransformer
SpecInfer w/ Incremental Decoding
SpecInfer w/ Sequence-based Speculative Inference
SpecInfer w/ Tree-based Speculative Inference
Per-token latency (ms)
0
5
10
15
20
25
30
35
40
BS=1
BS=2
BS=4
BS=8
BS=16
LLaMA-7B
(1 GPU/node, 1 node)
0
10
20
30
40
50
60
BS=1
BS=2
BS=4
BS=8
BS=16
OPT-30B
(4 GPUs/node, 1 node)
0
20
40
60
80
100
120
BS=1
BS=2
BS=4
BS=8
BS=16
LLaMA-65B
(4 GPUs/node, 2 nodes)