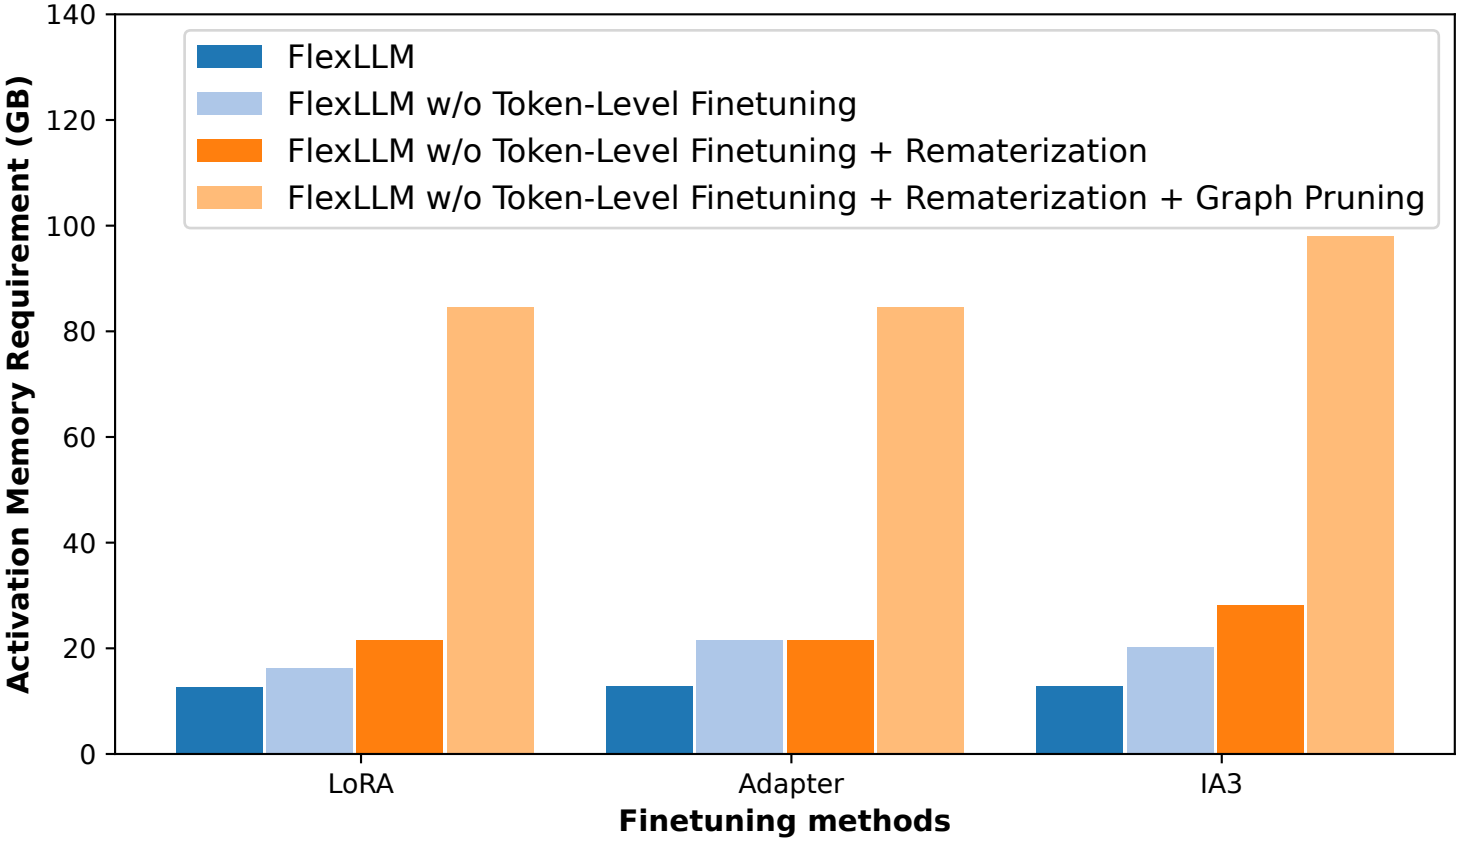
FlexLLM
FlexLLM w/o Token-Level Finetuning
FlexLLM w/o Token-Level Finetuning + Rematerization
FlexLLM w/o Token-Level Finetuning + Rematerization + Graph Pruning
Activation Memory Requirement (GB)
0
20
40
60
80
100
120
140
LoRA
Adapter
IA3
Finetuning methods