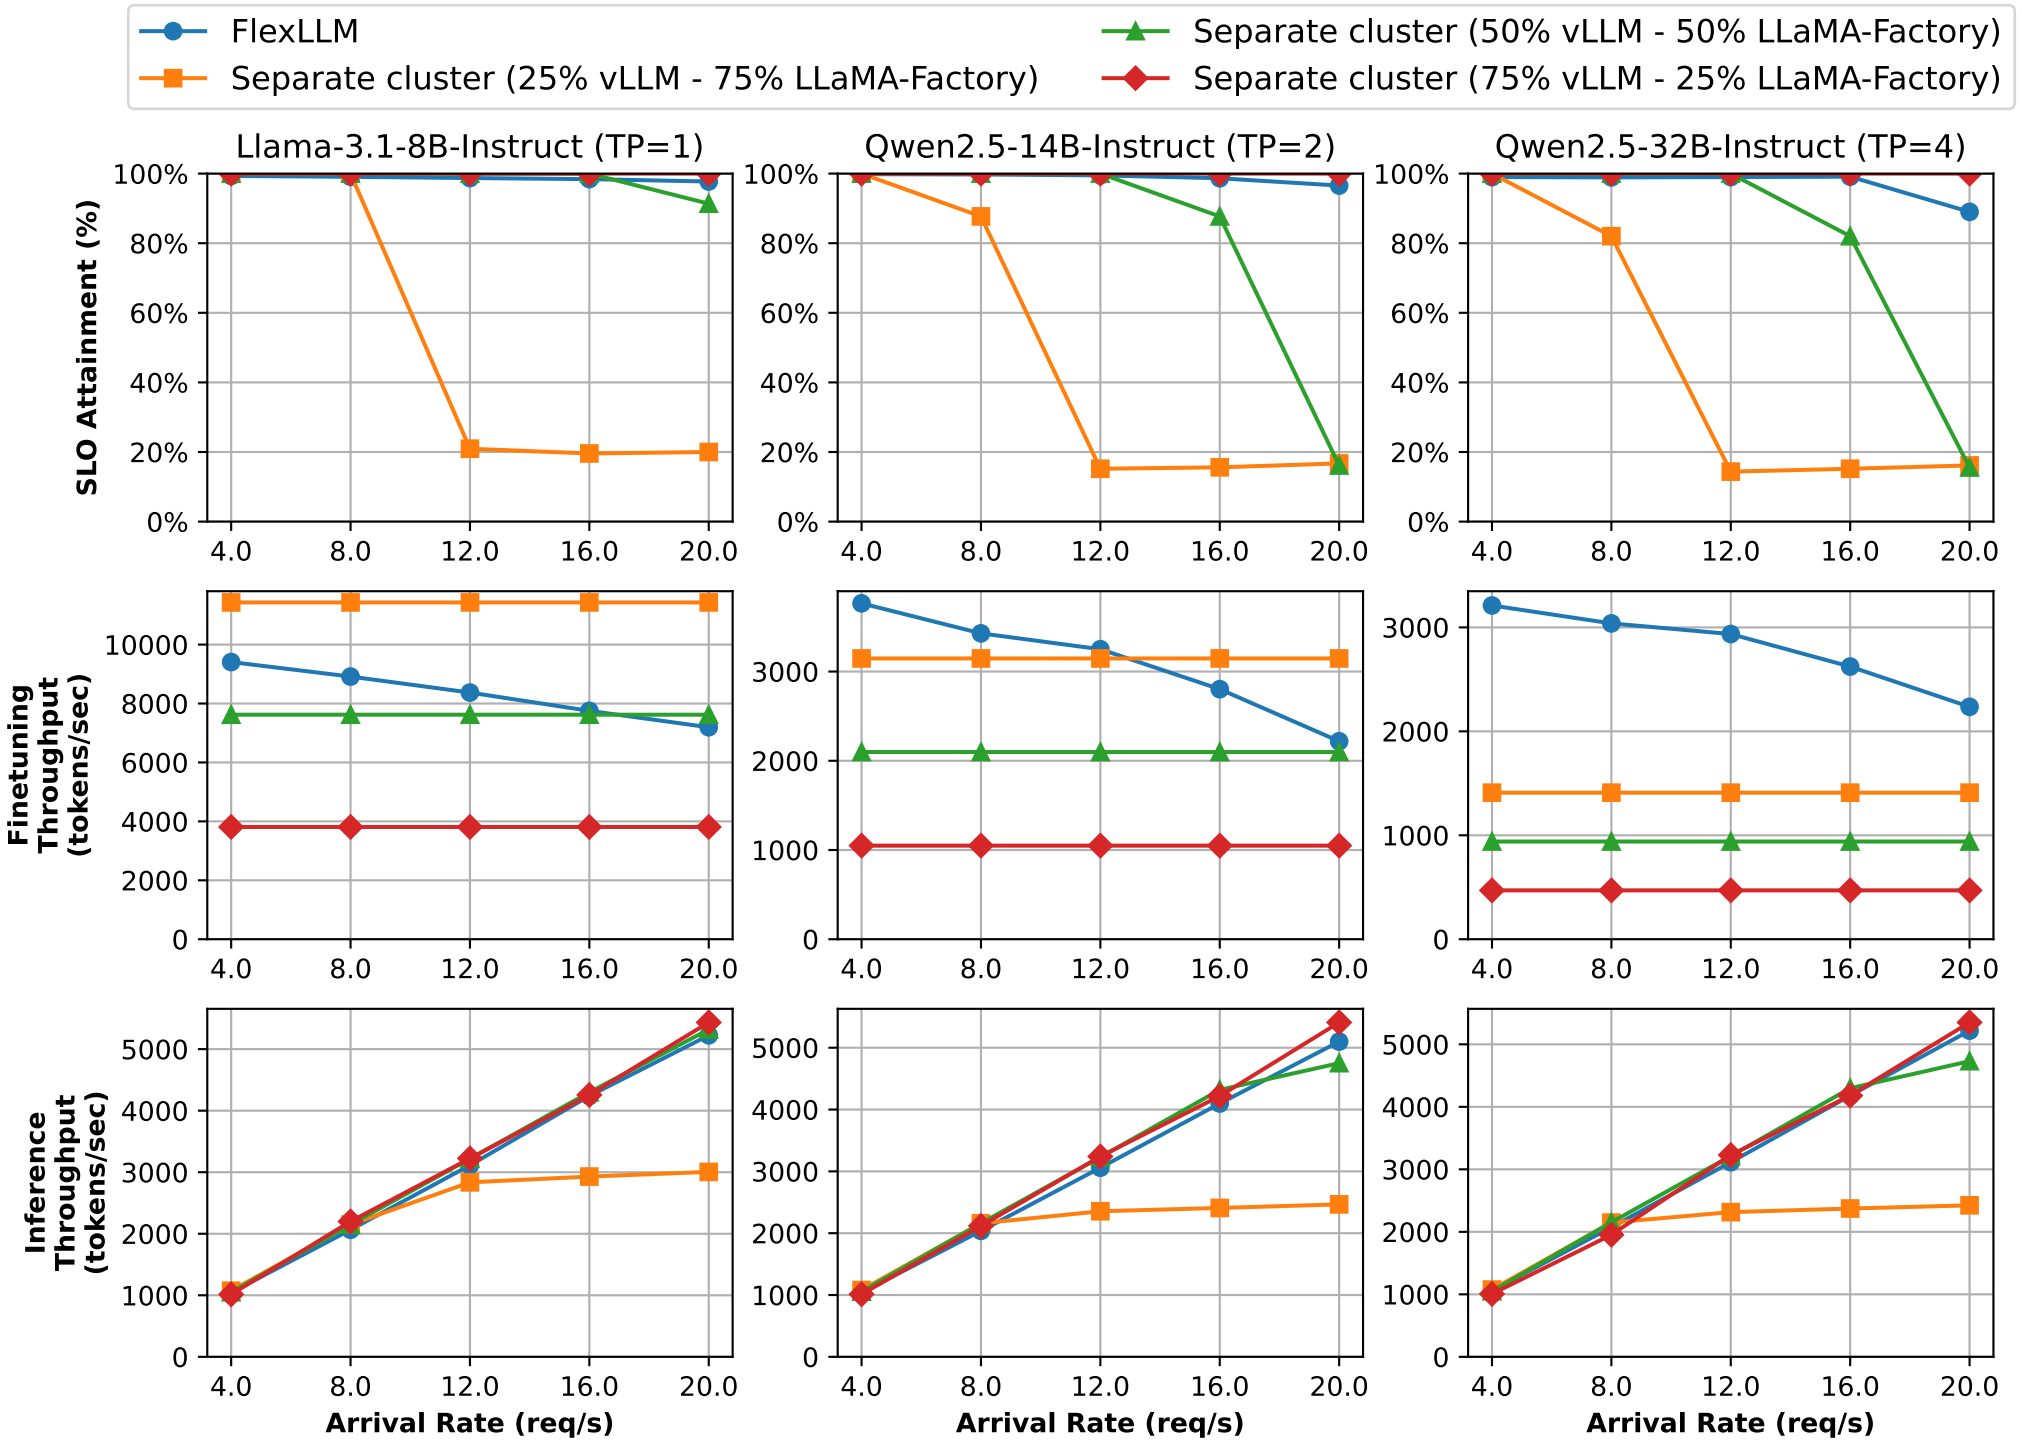
FlexLLM
Separate cluster (25% vLLM - 75% LLaMA-Factory)
Separate cluster (50% vLLM - 50% LLaMA-Factory)
Separate cluster (75% vLLM - 25% LLaMA-Factory)
Llama-3.1-8B-Instruct (TP=1)
Qwen2.5-14B-Instruct (TP=2)
Qwen2.5-32B-Instruct (TP=4)
SLO Attainment (%)
Finetuning Throughput (tokens/sec)
Inference Throughput (tokens/sec)
Arrival Rate (req/s)