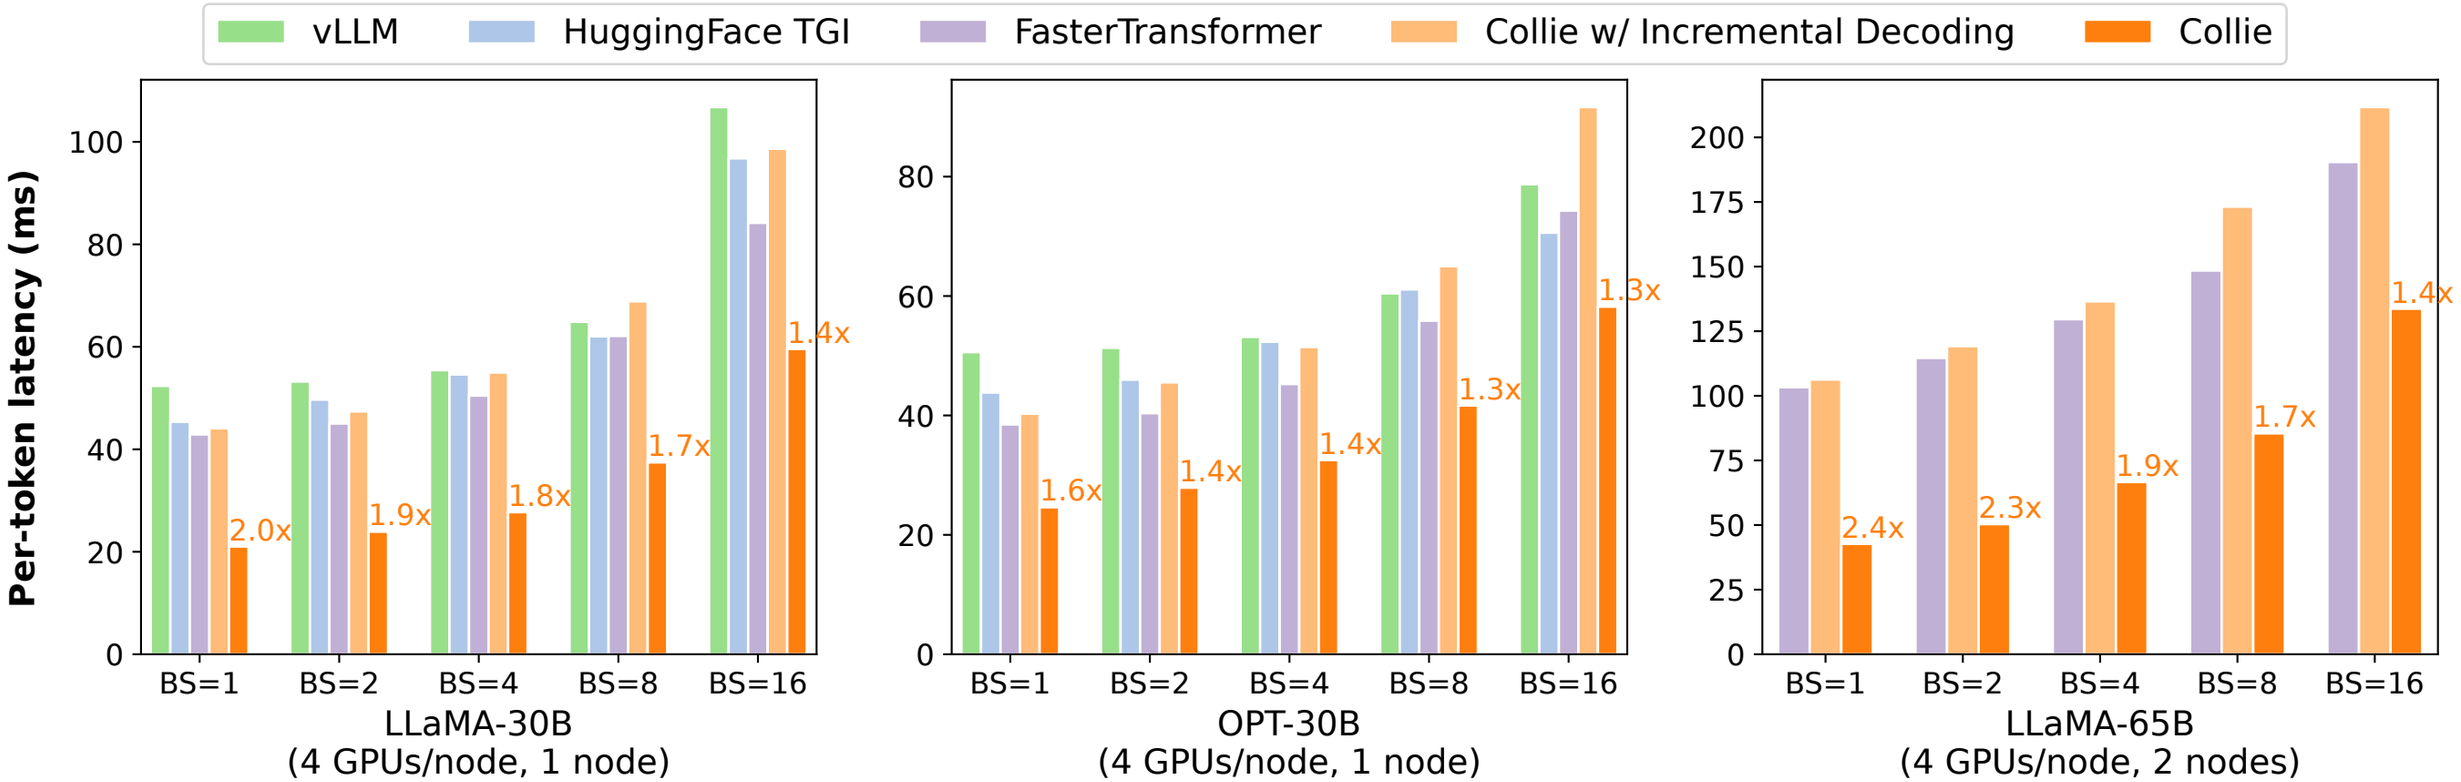
vLLM
HuggingFace TGI
FasterTransformer
Collie w/ Incremental Decoding
Collie
Per-token latency (ms)
0
20
40
60
80
100
2.0x
1.9x
1.8x
1.7x
1.4x
BS=1
BS=2
BS=4
BS=8
BS=16
LLaMA-30B
(4 GPUs/node, 1 node)
0
20
40
60
80
1.6x
1.4x
1.4x
1.3x
1.3x
BS=1
BS=2
BS=4
BS=8
BS=16
OPT-30B
(4 GPUs/node, 1 node)
0
25
50
75
100
125
150
175
200
2.4x
2.3x
1.9x
1.7x
1.4x
BS=1
BS=2
BS=4
BS=8
BS=16
LLaMA-65B
(4 GPUs/node, 2 nodes)