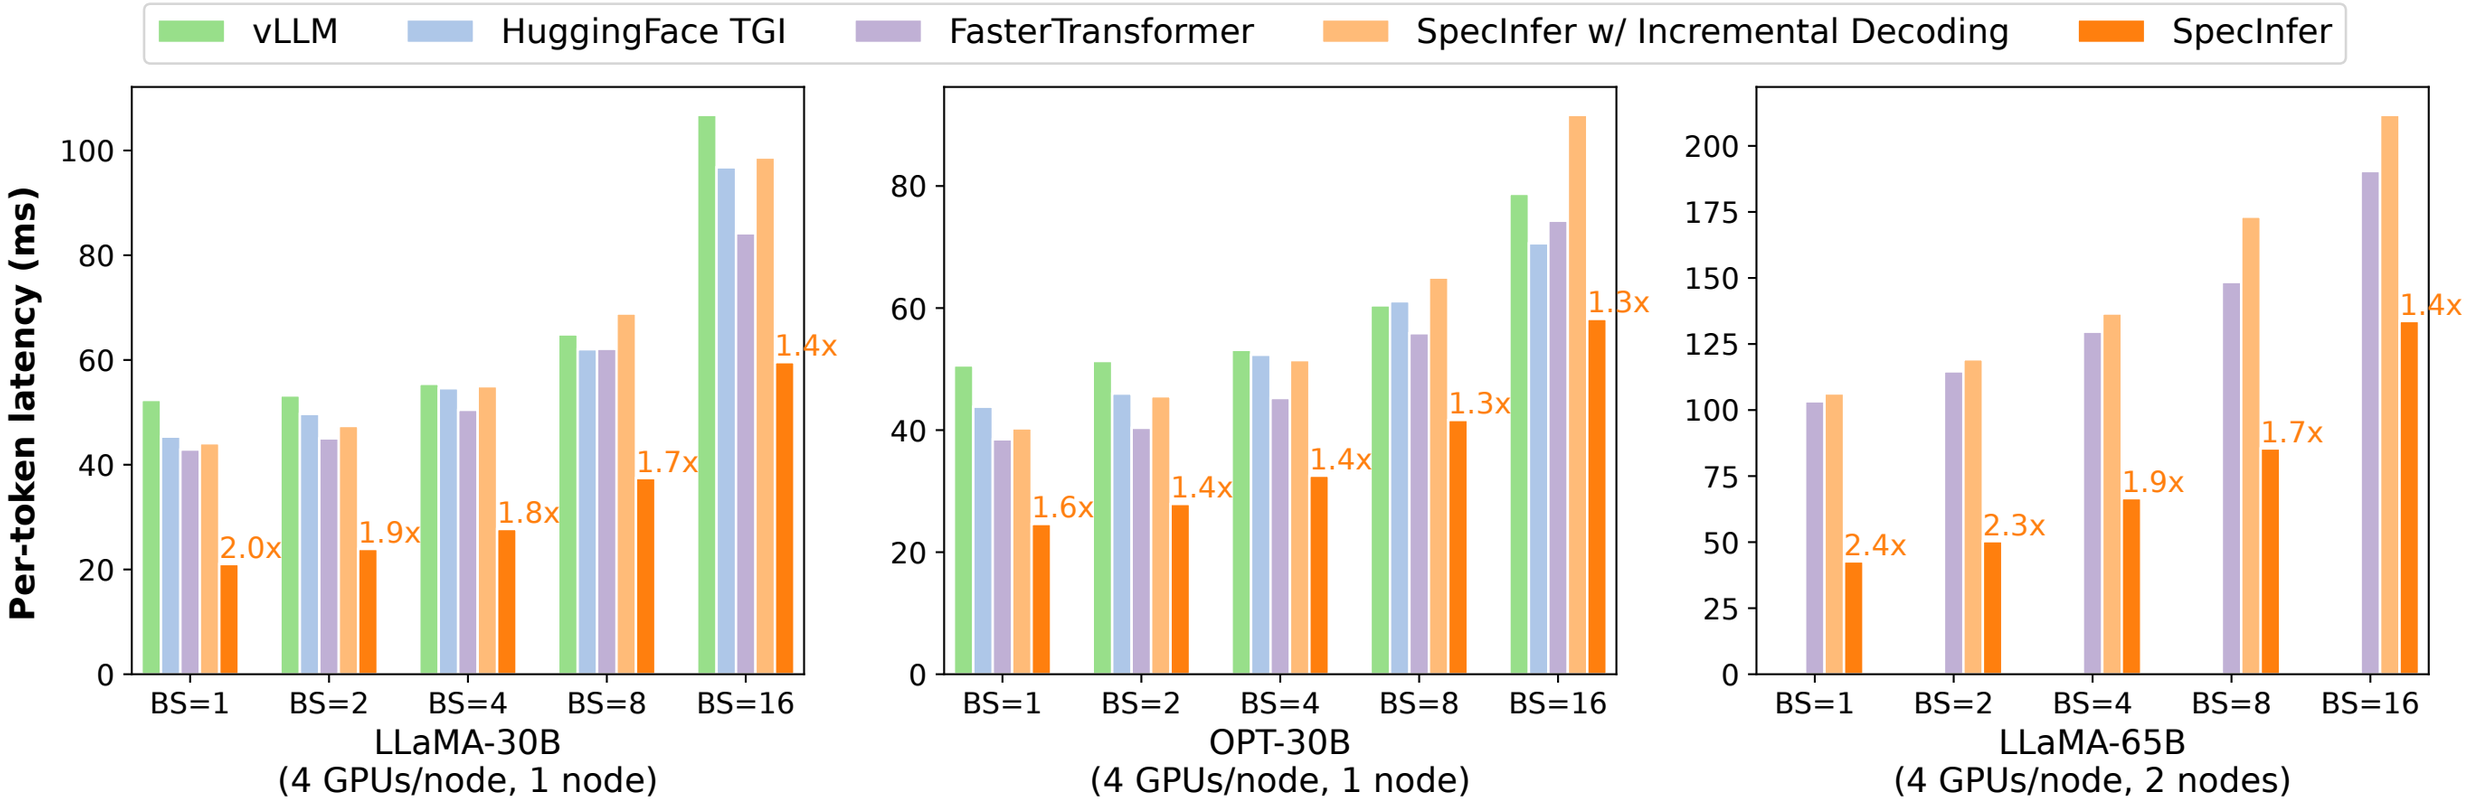
vLLM
HuggingFace TGI
FasterTransformer
SpecInfer w/ Incremental Decoding
SpecInfer
Per-token latency (ms)
BS=1
BS=2
BS=4
BS=8
BS=16
2.0x
1.9x
1.8x
1.7x
1.4x
LLaMA-30B
(4 GPUs/node, 1 node)
1.6x
1.4x
1.4x
1.3x
1.3x
OPT-30B
(4 GPUs/node, 1 node)
2.4x
2.3x
1.9x
1.7x
1.4x
LLaMA-65B
(4 GPUs/node, 2 nodes)